

In Asia the chat changes the delivery. In the West it only answers.

Conversational and agentic AI in parcel delivery. September 2026

Among the largest carriers in eight markets, **three let a conversational AI change a delivery**: Yamato on LINE since 2016, Japan Post and JD Logistics by voice. Three more do it through chat menus. All six are in Asia. **The Western carriers answer**, then send the customer to a web form. What decides whether an AI may change a delivery is not how clever the model is. It is **permission, data you can trust, and who is liable when it is wrong**.

01 What delivery companies have actually deployed

A conversational AI changes the delivery		3 operators
Yamato	Japan Post	JD Logistics
On LINE since June 2016. Change the date, time or place in conversation; the AI proposes the change itself	Voice AI rebooks a missed delivery (2020) and books pickups (2024), start to finish	Voice robot calls to book bulky deliveries and moves the time on request: 1.93m calls over 11.11 2021

A chat menu or script changes it		3 operators
SF Express	WeChat and app assistant since 2016: one-click address and time changes	
CJ Logistics	Chatbot menu includes address change; 22% of all inquiries go through it	
Delhivery	A scripted WhatsApp reply after a failed attempt books the re-attempt. Not AI	

Answers only, or a web form		the rest
DHL Freight (VIVA)	Chat, voice and callback for tracking and complaints; no booking claimed	
UPS	“AI-enabled intelligent assistants” in 20+ countries, scope unspecified; none named in its 2025 filing	
FedEx	Documented as unable to explain or act on a missed delivery	
PostNL, Australia Post, DHL Paket	Answer, then hand over to a person or a form	
La Poste, Swiss Post, USPS, Amazon	Changes happen on a web or app form	
Royal Mail, Evri, Correos	Nothing customer-facing found. Evri uses AI only for its own back-office staff	

Two much-quoted figures claim less than they seem: DHL's “93% answered without assistance” covers no booking, and UPS's “more than 98%” counts *humans and AI together*. **In the West an AI that changes a delivery is still almost unheard of; in East Asia it has run for a decade, inside the messaging app customers already use.** For a UAE operator that app is WhatsApp.

The real ceiling, and it is not the software

Three carriers publish the same limit in their own words. **UPS**: “Some shippers restrict the ability to change your delivery.” **FedEx**: shippers “may not allow recipients to change the destination once the shipment is in transit.” **Evri**: a parcel needing a signature cannot be redirected at all. What blocks a change of address is **the contract with whoever sent the parcel**, not the app the customer is using. An AI inherits every one of those locks.

02 What already works, with no AI at all

UPS My Choice launched in **October 2011** and had “more than 30 million members” by 2016. FedEx Delivery Manager

followed in **2013**, DHL On Demand Delivery after it. All three already let you change the day, use a pickup point, hold a parcel or redirect it, with no AI. Pricing follows cost: a **pickup point is free** because it saves the carrier a stop, while a **different home** costs \$5.55 within 120 miles at FedEx and \$14.50 to \$33.50 beyond. **AI does not change those economics.**

Serving yourself did not stop people getting in touch. Ofcom asked 4,058 UK adults in 2025: **68% had a delivery problem in six months**, and only 45% were satisfied with contacting the carrier.

Capgemini asked 2,874 shoppers and 500 executives what a delivery costs and what it is worth: **the company pays \$10.10**, the customer is **charged \$8.08**, and customers say it is **worth \$1.40**. Every delivery loses money before anything goes wrong; a failed one loses more.

03 What the AI vendors sell, and what it costs

None of these is a logistics company: they sell the same assistant to any industry. Carriers use them to answer questions; none of the Western carriers we checked has connected one to the shipment record so it can change a delivery.

Vendor, and what you are charged for	Price
Intercom (Fin) Only when the AI settles it	\$0.99
Yellow.ai Only when the AI settles it	\$0.99
Zendesk Only when a second AI agrees it was settled	\$1.50
Salesforce Agentforce Every conversation, settled or not	\$2.00
Cognigy Every conversation	\$0.72 \$0.10 at volume
Sierra, Decagon, Ada, PolyAI, Maven AGI No price published at all	

The headline prices cannot be compared. Intercom and Zendesk charge **only when it works**; Salesforce charges for **every conversation**. If the AI settles 7 in 10 by itself, Salesforce's \$2.00 really costs \$2.86 for each one settled, while Intercom's stays at \$0.99. Compare cost per problem solved. And paying only for success is less safe than it sounds: **the vendor decides what counts as success**, so ask for the right to audit and dispute it.

04 What fails

Ending a conversation is not solving a problem

The industry's favourite measure is *containment*: the share of conversations that finish without a human. It counts a customer who gave up as a win. Gartner predicts AI will settle **80% of common service issues by 2029**; it *measured* 14%

resolving on their own today (5,728 customers, 2024), and predicts **more than 40% of these projects will be cancelled by end-2027**.

Changing a record is the hard part, and it has been measured

There is a public test, τ -bench, that scores AI agents on **what the database looks like afterwards**, not on how pleasant the chat was. The best agents available finish **under 50%** of these tasks correctly, and run the same task eight times, **fewer than one in four** pass every time. The test was built by researchers tied to Sierra, an AI vendor.

Whatever the assistant says, the company said it

Air Canada's chatbot promised a grieving customer a bereavement discount its real policy did not allow. A Canadian tribunal held the airline to the bot's words: a chatbot is part of your website (*Moffatt*, 2024, CA\$812). A German appeal court went further in 2026: "the AI is a black box" was **rejected outright**, and a general "AI can make mistakes" disclaimer did not protect the company. The UAE has no AI law yet; a wrong delivery promise falls under ordinary consumer protection law.

Customer text plus permission to act is the way in for an attacker

Researchers hid instructions in an ordinary web form; Salesforce's AI later read them and obeyed (severity **9.4 out of 10**, September 2025). ServiceNow called a similar finding "works as intended". OWASP ranks this **the number one risk** for AI, with no reliable prevention; its advice is the **narrowest possible permissions** and a human sign-off for anything risky. A delivery note is the same kind of free text. DPD switched its chat AI off in January 2024 after a **routine software update** removed its safety rules.

05 What is different in the UAE

The channel is already WhatsApp, and an assistant is already on it

Emirates Post lists WhatsApp as its **only 24-hour channel** and calls it a "virtual assistant". Aramex made WhatsApp its main channel in 2018. FedEx put UAE delivery notifications on it in 2023. No official figure for WhatsApp use in the UAE exists; this rests on **what companies do**.

"The app itself is bad and unhelpful, and the WhatsApp service is also useless — it tells me there is no shipment linked to my data, although I actually have a shipment with Emirates Post."

Customer review, translated from Arabic

The assistant can talk; it cannot see the parcel. **The gap is data and permission.**

What UAE customers actually complain about

In 454 negative app-store reviews of UAE delivery companies, the biggest complaint is **not being able to reach a human**, at 22.0%. Split by company, it flips over:

Share of complaints	This operator	Other UAE carriers
Something they were not allowed to do Move a delivery, change an address	17.2%	10.9%
Could not reach a human	9.3%	28.4%

This company's customers complain about things they were not allowed to do. Everyone else's complain about not reaching anyone. That flip is the argument for what to build. One finding cuts the other way: 6.6% of complaints describe a delivery attempt the customer says never happened, matching the contradictory rows in the supplied data exactly. An assistant that reads “failed - customer not available” back to those people will make them angrier than no assistant at all.

Two costs that apply whoever builds it
WhatsApp stops being free for service replies on 1 October 2026. In the UAE each message costs about \$0.0157. The prototype's conversations send about three replies, so roughly \$0.04 each: three times what the AI model itself cost (\$0.013, measured). It applies whether you buy or build.

Recommendation: buy the platform, build the rules
Buy the messaging channels, integrations, testing and reporting: the same for everyone, and open standards now make changing vendor much easier.
Build the part that is yours: the actions that change your records, and the rules deciding when they are allowed, worked out from the shipment file rather than written into the model's instructions. They depend on facts only you hold (cash owed, attempts made, records that disagree), and you are liable for them (section 04). Hand them to the platform through a documented interface; they stay in your code, under your own tests.

Gulf Arabic is the weakest kind of Arabic for these models, not the strongest. Every model tests better on formal written Arabic than on spoken dialects, Emirati included. A 2026 study of ~30,000 *delivery* support messages found that testing with translated text **flatters the results**. The loudest vendor claims “95%+ dialect recognition” and publishes no test set or method.

06 Buy or build

Per 100 conversations, if the AI settles 7 in 10	Cost
Build: the AI model's running cost, measured on the prototype	\$1.30
Buy from Intercom, \$0.99 each time it works	\$69
Buy from Zendesk, \$1.50 each time it works	\$105
Buy from Agentforce, \$2.00 per conversation	\$200

Running the model is under 2% of what a platform charges; the rest pays for channels, integrations, testing, handover to a human, reporting and, when you pay per success, someone else carrying the risk. Menlo Ventures found companies building their own AI dropped from 47% to 24% in a single year.

Four numbers everyone quotes, that I deliberately did not use

- \$17.78 per failed delivery.** A 2020 vendor press release based on a survey of 1,200 American *shoppers*, who cannot know what a delivery costs a retailer. The most-copied figure in this field.
- 99% of Aramex enquiries handled by the bot.** One vendor's undated marketing page. A different vendor says 50% for the same bot.
- “Where is my order” is 25 to 40% of all contacts, at £3 to £5 each.** Every source I found is a vendor blog quoting another vendor blog.
- 1 in 10 parcels fail first time.** A misreading: the original counts *companies* that see that rate, not parcels. IMRG's measurement is 2.77%.

Sources. Every figure has a named source and a date. The full list, with each source's type and the claims that failed checking, is in *research/* in the repository. **Operators covered:** the largest carriers in the US, Europe, Australia, China, Japan, Korea, India and the Gulf, each searched in its own language; not Latin America or Africa. **Reviews:** 718 collected (January 2025 to September 2026), the 454 negative ones sorted by hand; a re-check of 40 agreed with the labels ~82% of the time. **Limits.** App-store reviews show where people complain, not satisfaction or call volumes, and Arabic is 13.7% of the sample. No claim about either is drawn from them.