

Plenty of delivery companies have an AI that **answers questions**. Across the whole public record, **exactly one** has an AI that **changes a delivery**: one that moves the date or the address and writes it back to the record. The thing stopping the rest is not how clever the model is. It is **permission, data you can trust, and who is liable when it is wrong**.

The brief asks four questions. Sections 01 and 02 cover what works today, 03 what it costs, 04 what fails, 05 what is different in the UAE, and 06 what to buy and what to build.

01 What delivery companies have actually deployed

Operator	Can it change a delivery?	What is on the record
Japan Post	Yes	A phone AI rebooks a missed delivery (Nov 2020) and arranges pickups (Feb 2024), start to finish
DHL Freight (VIVA)	No, answers only	Tracking and complaint help by chat, voice and callback. It claims no booking or rescheduling
UPS	Never says	“AI-enabled intelligent assistants” in 20+ countries, scope unspecified. Its 2025 annual filing names no assistant at all
FedEx	No, answers only	Its assistant is documented as unable to explain a missed delivery or do anything about it
PostNL, Australia Post	No, answers only	Answer, then pass to a person. Australia Post says 55% of enquiries are “supported by self-serve AI options”
Royal Mail, Evri, Correos	No AI at all	Nothing customer-facing found. Evri uses AI only for its own back-office staff

Japan Post's assistant collects three things by phone (tracking number, the day you want, a contact number) and then rebooks the delivery. Two numbers here get quoted a lot and both need care. DHL's “**93%** of enquiries answered without assistance” is from DHL's own marketing site and covers no booking. UPS's “more than **98%**” counts requests its *humans and AI together* touch, not what the AI settles alone. **An AI that changes a delivery is almost unheard of at a Western carrier**. That is the opening, and the reason to doubt anyone selling it as a solved category.

The real ceiling, and it is not the software

Three carriers publish the same limit in their own words. **UPS**: “Some shippers restrict the ability to change your delivery.” **FedEx**: shippers “may not allow recipients to change the destination once the shipment is in transit.” **Evri**: a parcel that needs a signature cannot be redirected at all. What blocks a change of address is **the contract with whoever sent the parcel**, not the app the customer is using. An AI inherits every one of those locks. Any pitch suggesting AI can do something the website cannot already do is simply wrong.

02 What already works, with no AI at all

UPS My Choice launched in **October 2011** and had “more than 30 million members” by 2016. FedEx Delivery Manager launched in **2013**, DHL On Demand Delivery after it. All three already let you change the day, send a parcel to a pickup point, hold it or redirect it. No AI, for over a decade. The pricing is the same everywhere and has nothing to do with intelligence: a **pickup point is free** because it saves the carrier a stop, while **a different home** is where FedEx charges **\$5.55** within 120 miles and **\$14.50 to \$33.50** beyond. **AI sits on top of those economics. It does not change them**.

And serving yourself did not stop people getting in touch. Ofcom asked **4,058** UK adults in 2025: **68% had a delivery problem in six months**, and only **45%** were satisfied with contacting the carrier.

The money was already tight before any of this. Capgemini asked **2,874** shoppers and **500** executives what a delivery costs and what it is worth: **the company pays \$10.10**, the customer is **charged \$8.08**, and customers say it is **worth \$1.40**. Every delivery loses money before anything goes wrong, which is why a failed one hurts and why nobody can charge their way out of it.

03 What the AI vendors sell, and what it costs

Vendor	What you are charged for	Price
Intercom (Fin)	Only when the AI settles it	\$0.99
Yellow.ai	Only when the AI settles it	\$0.99
Zendesk	Only when a second AI agrees it was settled	\$1.50
Salesforce Agentforce	Every conversation, settled or not	\$2.00
Cognigy	Every conversation	\$0.72, or \$0.10 at volume
Sierra, Decagon, Ada, PolyAI, Maven AGI	No price published at all	

The headline prices cannot be compared, and that is the most useful thing to know before talking to any of them. Intercom and Zendesk charge **only when it works**. Salesforce charges for **every conversation**. Suppose the AI settles 7 in 10 by itself: Salesforce's **\$2.00** then really costs **\$2.86** for each one it settled, while Intercom's stays at **\$0.99**. Convert every quote into cost per problem actually solved before comparing. And paying only for success is less safe than it sounds, because **the vendor also decides what counts as success**. Ask for the right to audit it, and a way to dispute it.

04 What fails

Ending a conversation is not solving a problem

The industry's favourite measure is *containment*: the share of conversations that finish without a human. It counts a customer who gave up as a win. Gartner predicts AI will settle **80% of common service issues by 2029**. Gartner also *measured* how many people get an issue resolved on their own today: **14%** (5,728 customers, 2024). And it predicts **more than 40% of these projects will be cancelled by end-2027** for cost, unclear value and weak controls. All three numbers are theirs, and the distance between them is the risk.

Changing a record is the hard part, and it has been measured

There is a public test, τ -*bench*, that scores AI agents on **what the database looks like afterwards**, not on how pleasant the chat was. That is exactly what a delivery assistant does. The best agents available finish **under 50%** of these tasks correctly. Run the same task eight times and **fewer than one in four** pass every single time. The test was built by researchers tied to Sierra, an AI vendor, which makes the result harder to dismiss, not easier.

Whatever the assistant says, the company said it

Air Canada's chatbot told a grieving customer he could claim a bereavement discount after buying his ticket. The airline's real policy said otherwise. A Canadian tribunal held the company to what the bot said: a chatbot is part of your website, and having the correct policy on another page is no defence (*Moffatt*, 2024, worth **CA\$812**, a small case that became the reference point). A German appeal court went further in 2026: “the AI is a black box” was **rejected outright**, and a general “AI can make mistakes” disclaimer did not protect the company. The UAE has no AI law yet, so a wrong delivery promise would be judged under ordinary consumer protection law, which does not care that a machine said it.

Customer text plus permission to act is the way in for an attacker

Researchers hid instructions inside an ordinary web form. Salesforce's AI read them later and obeyed. The flaw was rated **9.4 out of 10** for severity, September 2025. A similar finding against ServiceNow got the reply that the system “**works as intended**”. OWASP, which publishes the standard security list, ranks this **the number one risk** for AI and admits there is no reliable way to prevent it; its advice is the **narrowest possible permissions** plus a human sign-off for anything risky. A delivery note saying “leave at reception” is the same kind of free text as that form. It rarely goes wrong, but publicly when it does: DPD switched its chat AI off in January 2024 after a **routine software update** quietly removed the safety rules.

05 What is different in the UAE

The channel is already WhatsApp, and an assistant is already on it

If this operator looks like a UAE postal operator, the channel question is already answered. Emirates Post lists WhatsApp as its **only 24-hour channel** and calls it a “virtual assistant”. Aramex made WhatsApp its main channel in 2018. FedEx put UAE delivery notifications on it in 2023. No official figure for WhatsApp use in the UAE appears to exist, so this rests on **what companies do, not on a statistic**.

“The app itself is bad and unhelpful, and the WhatsApp service is also useless — it tells me there is no shipment linked to my data, although I actually have a shipment with Emirates Post.”

Customer review, translated from Arabic

That is not a language problem. The assistant can talk; it cannot see the parcel. **The gap is data and permission, which is the whole argument of this document.**

What UAE customers actually complain about

I collected **718** public app-store reviews of UAE delivery companies (January 2025 to September 2026) and sorted the **454** negative ones by hand, using rules written after reading all of them. I re-checked 40 of them

myself and agreed with the labels ~82% of the time. Across the whole set the biggest complaint is **not being able to reach a human**, at 22.0%. Split by company, it flips over:

What they complain about	This operator	Other UAE carriers
Something they could not do (move a delivery, change an address)	17.2%	10.9%
Could not reach a human	9.3%	28.4%

This company's customers complain about things they were not allowed to do. Everyone else's complain about not reaching anyone. That flip is the argument for what to build. One finding cuts the other way and belongs on the record: 6.6% of complaints describe a delivery attempt the customer says never happened, matching the contradictory rows in the supplied data exactly. An assistant that reads “failed - customer not available” back to those people will make them angrier than no assistant at all.

Two costs that apply whoever builds it

WhatsApp stops being free for service replies on 1 October 2026. In the UAE each message costs about \$0.0157, so a conversation of four to eight replies adds roughly \$0.06 to \$0.13. That is about the same as running the AI model itself, and it applies whether you buy or build.

Gulf Arabic is the weakest kind of Arabic for these models, not the strongest. The public tests show every model doing better on formal written Arabic than on the dialects people actually speak, Emirati included. Worse for a buyer, a 2026 study of ~30,000 anonymised delivery support messages found that testing with translated text flatters the results. Among vendors selling Arabic the loudest claims “95%+ dialect recognition” and publishes no test set or method. Nobody publishes an Arabic number you can check.

06 Buy or build

Per 100 conversations, if the AI settles 7 in 10	Cost
Build: the AI model's own running cost only	\$10
Buy from Intercom, \$0.99 each time it works	\$69
Buy from Zendesk, \$1.50 each time it works	\$105
Buy from Agentforce, \$2.00 per conversation	\$200

Running the model is only 3 to 15% of what a platform charges. Anyone saying “build it, the AI is cheap” is answering a question nobody asked. The other 90% is the messaging channels, the ready-made links into other systems, the testing, the handover to a human with the conversation attached, the reporting, and when you pay per success, someone else carrying the risk. Menlo Ventures found companies building their own AI dropped from 47% to 24% in a single year.

Recommendation: buy the platform, build the rules

Buy the messaging channels, the links into other systems, the testing and the reporting. It is the same for everyone, it is where 90% of the money goes, and open standards now make changing vendor much easier than it was.

Build the small part that is actually yours: the actions that change your own records, and the rules deciding when they are allowed. Work those rules out from the shipment file rather than writing them into instructions for the model. They are what you are liable for (section 04) and they depend on facts only you hold: money owed on delivery, attempts already made, whether two records for one parcel disagree. Bought agents are built for the average company. Hand the rules over through a documented interface and they stay in your code, under your own tests.

Four numbers everyone quotes, that I deliberately did not use

- **\$17.78 per failed delivery.** A vendor press release from 2020, based on a survey of 1,200 American shoppers, which cannot tell you what a delivery costs a retailer. The most-copied figure in this field; a business case built on it falls over at the first question.
- **99% of Aramex enquiries handled by the bot.** One vendor's undated marketing page. A different vendor says 50% for the same bot.
- **“Where is my order” is 25 to 40% of all contacts, at £3 to £5 each.** Every source I found is a vendor blog quoting another vendor blog.
- **1 in 10 parcels fail first time.** A misreading. The original says what share of companies see that rate, not what share of parcels fail. The best real measurement, by the UK retail trade body IMRG, is 2.77%.

Sources. Every figure here has a named source and a date. The full list is in research/ in the accompanying repository, with the scripts that collected and sorted the reviews. It carries the web addresses, a label on each source saying whether it is a company filing, a regulator, journalism or vendor marketing, and a register of claims that did not survive checking. Limits of my own research. App-store reviews are a place people complain, not a measure of satisfaction, and say nothing about how many calls a company gets. Arabic is under-represented at 13.7% of the sample. No claim about call volumes or customer satisfaction can be drawn from it, and none is made.